

Title: Predicting Customer Churn in a Telecommunications Company Objective: Develop a predictive model to accurately identify telecom customers who are likely to churn, enabling the company to take proactive measures to retain them.

Business Context: Customer churn is a significant issue for telecommunications companies, leading to substantial revenue loss. Understanding and predicting customer churn is critical for developing effective retention strategies. By analyzing customer data, we aim to identify the key factors contributing to churn and build a model that can predict at-risk customers.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: # importing the data.
data=pd.read_csv(r"C:\Users\medam\Downloads\archive (8)\WA_Fn-UseC_-Telco-Customer-Churn
# Head gives the top 5 records.
data.head()
```

```
Out[2]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic

5 rows × 21 columns

```
In [3]: # checking the dimensions of the data.
data.shape
```

```
Out[3]: (7043, 21)
```

```
In [4]: # checking the types of columns present in the data.
data.columns
```

```
Out[4]: Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
              'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
              'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
              'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',
              'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
              dtype='object')
```

```
In [5]: # checking the data type of each column.
data.dtypes
```

```
Out[5]: customerID      object
gender          object
SeniorCitizen   int64
Partner         object
Dependents      object
tenure          int64
PhoneService    object
MultipleLines   object
InternetService object
OnlineSecurity  object
OnlineBackup    object
DeviceProtection object
TechSupport     object
StreamingTV     object
StreamingMovies object
Contract        object
PaperlessBilling object
PaymentMethod   object
MonthlyCharges  float64
TotalCharges    object
Churn           object
dtype: object
```

```
In [6]: # checking the descriptive statistics of numerical variables.
data.describe()
```


```
Out[6]:
```

	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

- senior citizen supposed to be a categorical variable that is why the 25%-50%-75% distribution is not proper.
- 75% customers have tennure less than 55 months.
- The average monthly charges are USD65 but the customers are paying USD89.

```
In [7]: data["Churn"].value_counts().plot(kind="bar",figsize=(8,6))
plt.xlabel("churn")
plt.ylabel("target variable",labelpad=14)
plt.title("count of target variable per category")
```

```
Out[7]: Text(0.5, 1.0, 'count of target variable per category')
```



```
In [8]: data["Churn"].value_counts()
```

```
Out[8]: No      5174  
       Yes      1869  
       Name: Churn, dtype: int64
```

```
In [9]: # checking the percentage of distribution in churn.  
       data["Churn"].value_counts(normalize=True)*100
```

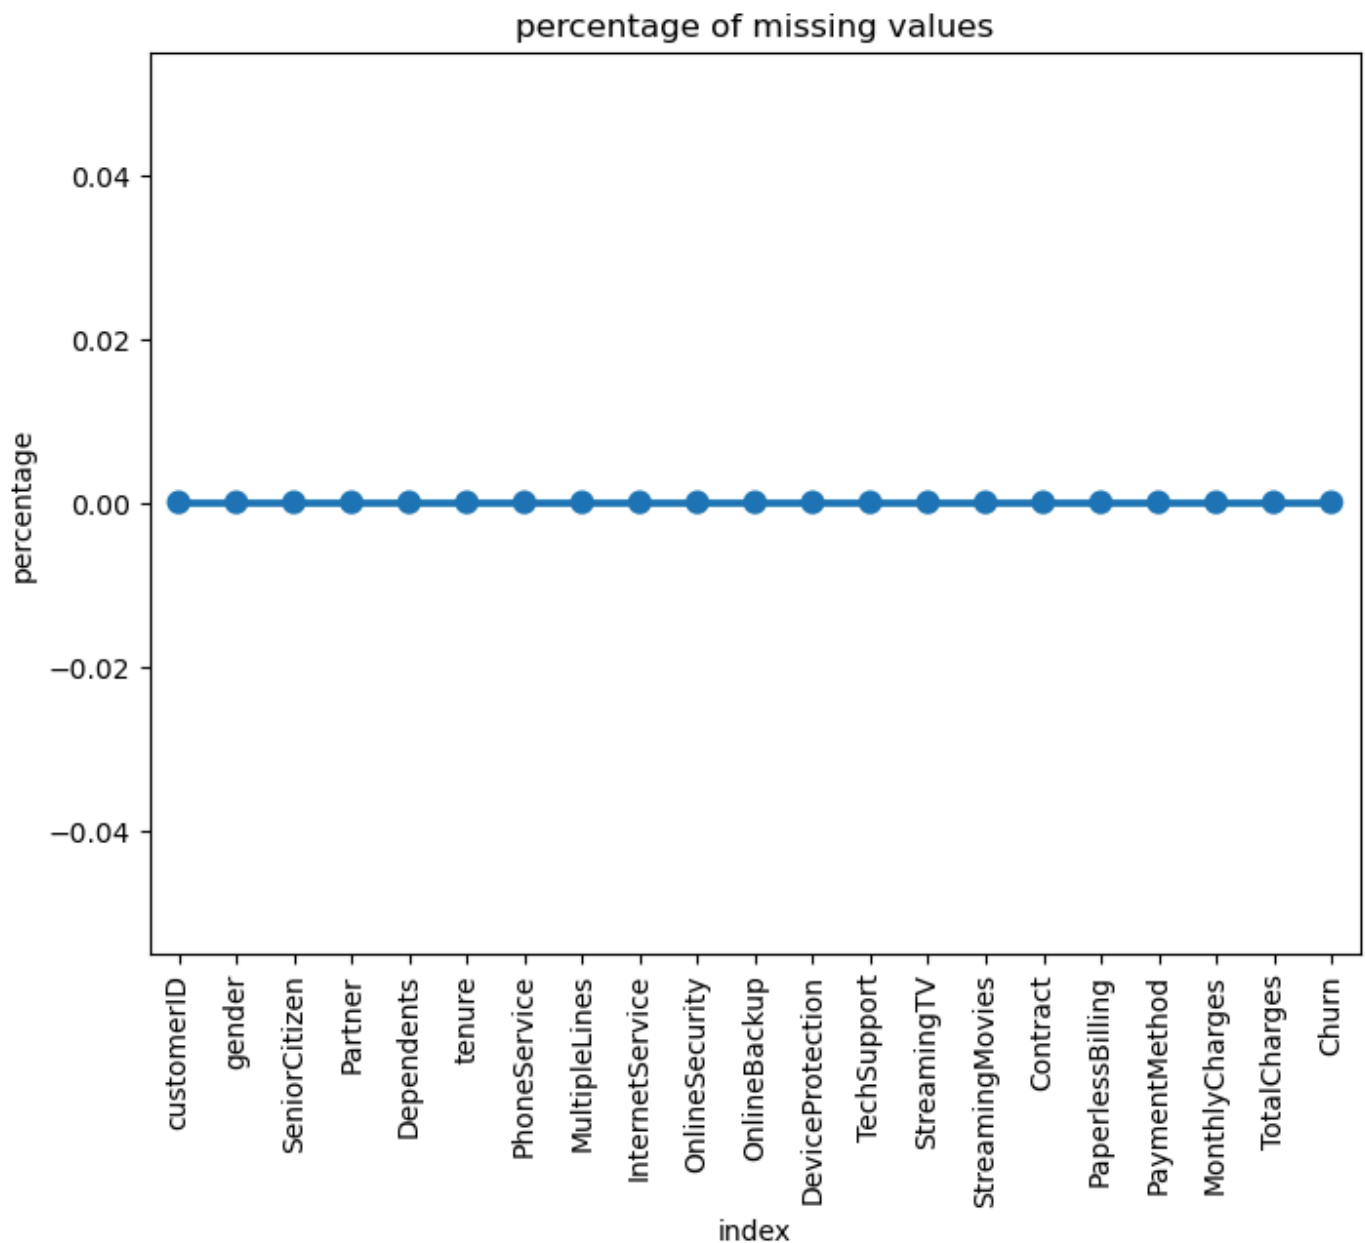
```
Out[9]: No      73.463013  
       Yes      26.536987  
       Name: Churn, dtype: float64
```

- Here we can see the data is imbalanced(73:27) ratio.
- so we analyze the data with other features while taking the target values separately to get some insights.

```
In [10]: # Here we are using verbose is True because we have many columns.  
       data.info(verbose=True)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                7043 non-null   object
2   SeniorCitizen         7043 non-null   int64
3   Partner               7043 non-null   object
4   Dependents            7043 non-null   object
5   tenure               7043 non-null   int64
6   PhoneService          7043 non-null   object
7   MultipleLines         7043 non-null   object
8   InternetService       7043 non-null   object
9   OnlineSecurity        7043 non-null   object
10  OnlineBackup          7043 non-null   object
11  DeviceProtection      7043 non-null   object
12  TechSupport           7043 non-null   object
13  StreamingTV           7043 non-null   object
14  StreamingMovies       7043 non-null   object
15  Contract              7043 non-null   object
16  PaperlessBilling      7043 non-null   object
17  PaymentMethod         7043 non-null   object
18  MonthlyCharges        7043 non-null   float64
19  TotalCharges          7043 non-null   object
20  Churn                 7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
In [11]: # To identify missing the percentage of values.
missing=pd.DataFrame((data.isnull().sum()*100/data.shape[0]).reset_index())
plt.figure(figsize=(8,6),dpi=100)
ax=sns.pointplot(x="index",y=0,data=missing)
plt.xticks(rotation=90)
plt.ylabel("percentage")
plt.title("percentage of missing values")
plt.show()
```



Missing data- initial intuition.

- Here we can see there is no missing data in this dataset. ### General thumb rules.
- If the variable has lower no.of missing values then we can use mean/median/mode(it depends on type of the variable).If the variable is numerical we can use mean/median whereas if the data is categorical we can use mode.
- If the variable has higher no.of missing values(60-70%) then undoubtedly we can drop that variable.

Data cleaning.

- creating a copy of base data for manipulation and processing.

```
In [12]: data=data.copy()
```

- Total charges supposed to be numerical.so we will convert that into numerical datatype.

```
In [13]: data.TotalCharges=pd.to_numeric(data.TotalCharges,errors="coerce")
data.isnull().sum()
```

```
Out[13]: customerID      0
gender      0
SeniorCitizen  0
Partner      0
Dependents    0
tenure      0
PhoneService  0
MultipleLines  0
InternetService  0
OnlineSecurity  0
OnlineBackup  0
DeviceProtection  0
TechSupport    0
StreamingTV    0
StreamingMovies  0
Contract      0
PaperlessBilling  0
PaymentMethod  0
MonthlyCharges  0
TotalCharges   11
Churn          0
dtype: int64
```

- we can see some of the missing values in total charges.so,let's see.

```
In [14]: data.loc[data["TotalCharges"].isnull()==True]
```

Out[14]:	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetServ
488	4472-LVYGI	Female	0	Yes	Yes	0	No	No phone service	I
753	3115-CZMZD	Male	0	No	Yes	0	Yes	No	
936	5709-LVOEQ	Female	0	Yes	Yes	0	Yes	No	I
1082	4367-NUYAO	Male	0	Yes	Yes	0	Yes	Yes	
1340	1371-DWPAZ	Female	0	Yes	Yes	0	No	No phone service	I
3331	7644-OMVMY	Male	0	Yes	Yes	0	Yes	No	
3826	3213-VVOLG	Male	0	Yes	Yes	0	Yes	Yes	
4380	2520-SGTTA	Female	0	Yes	Yes	0	Yes	No	
5218	2923-ARZLG	Male	0	Yes	Yes	0	Yes	No	
6670	4075-WKNIU	Female	0	Yes	Yes	0	Yes	Yes	I
6754	2775-SEFEE	Male	0	No	Yes	0	Yes	Yes	I

11 rows × 21 columns

Missing value treatment.

- since, the missing records are very low compared to total dataset is very low.so,it is safe to ignorr them from further processing.

```
In [15]: data.dropna(how="any", inplace=True)
```

```
In [16]: data.shape
```

```
Out[16]: (7032, 21)
```

- Divide customers into bins based on tennure.

```
In [17]: data["tenure"].max()
```

```
Out[17]: 72
```

```
In [18]: data["tenure"].min()
```

```
Out[18]: 1
```

```
In [19]: # Group the tennure in bins of 12 months.
labels=["{0} - {1}".format(i,i+11) for i in range(1,72,12)]
data["tenure_group"]=pd.cut(data["tenure"],range(1,80,12),right=False,labels=labels)
data["tenure_group"].value_counts()
```

```
Out[19]: 1 - 12      2175
61 - 72      1407
13 - 24      1024
25 - 36       832
49 - 60       832
37 - 48       762
Name: tenure_group, dtype: int64
```

- Remove the columns which are not required.

```
In [20]: data.drop(columns=["customerID", "tenure"], inplace=True)
```

```
In [21]: data.head()
```


```
Out[21]:
```



	gender	SeniorCitizen	Partner	Dependents	PhoneService	MultipleLines	InternetService	OnlineSecurity	Onl
0	Female	0	Yes	No	No	No phone service	DSL	No	
1	Male	0	No	No	Yes	No	DSL	Yes	
2	Male	0	No	No	Yes	No	DSL	Yes	
3	Male	0	No	No	No	No phone service	DSL	Yes	
4	Female	0	No	No	Yes	No	Fiber optic	No	

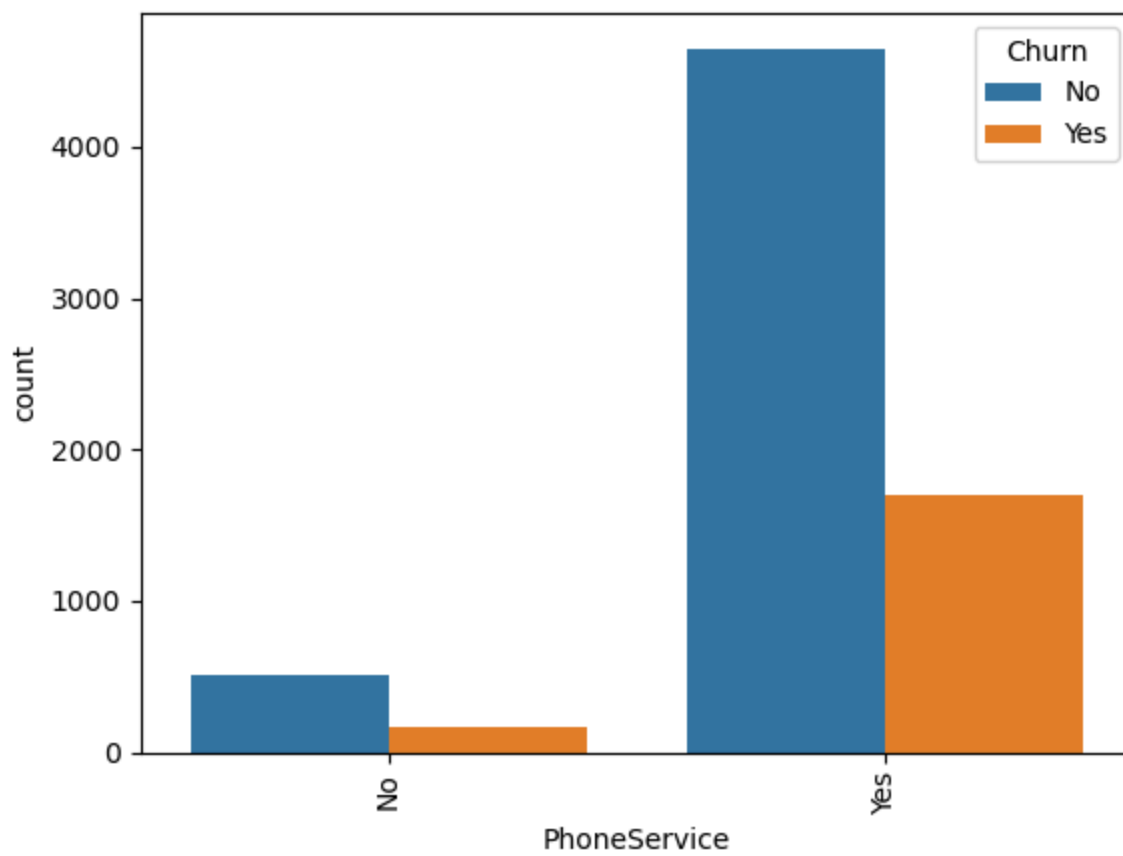

- plot distribution of individual predictors by churn.



Univariate Analysis.

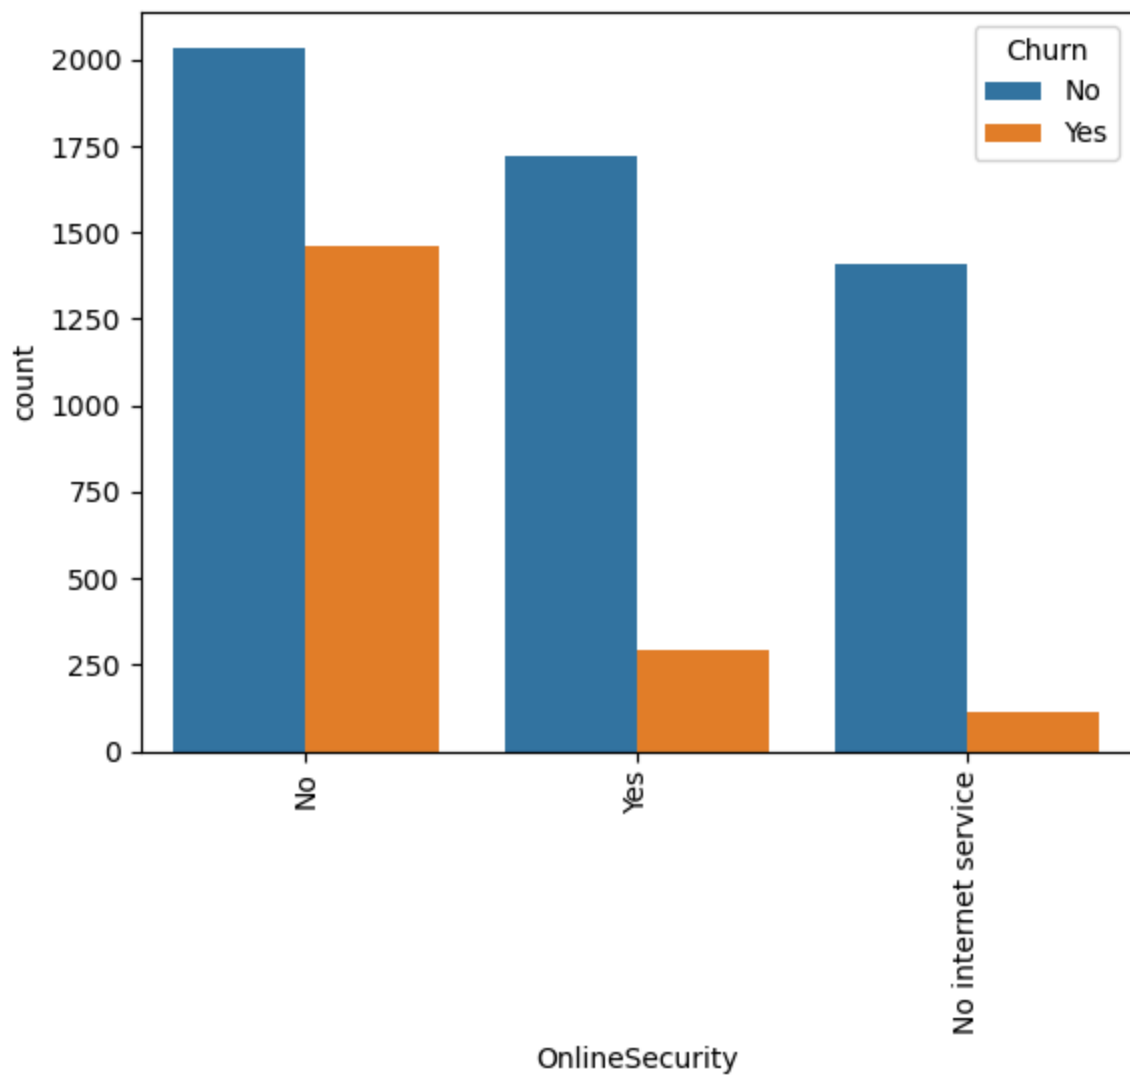
```
In [22]: for i, predictor in enumerate(data.drop(columns=["MonthlyCharges", "TotalCharges", "Churn"]):  
    plt.figure(i)  
    sns.countplot(data=data, x=predictor, hue="Churn")  
    plt.xticks(rotation=90)
```







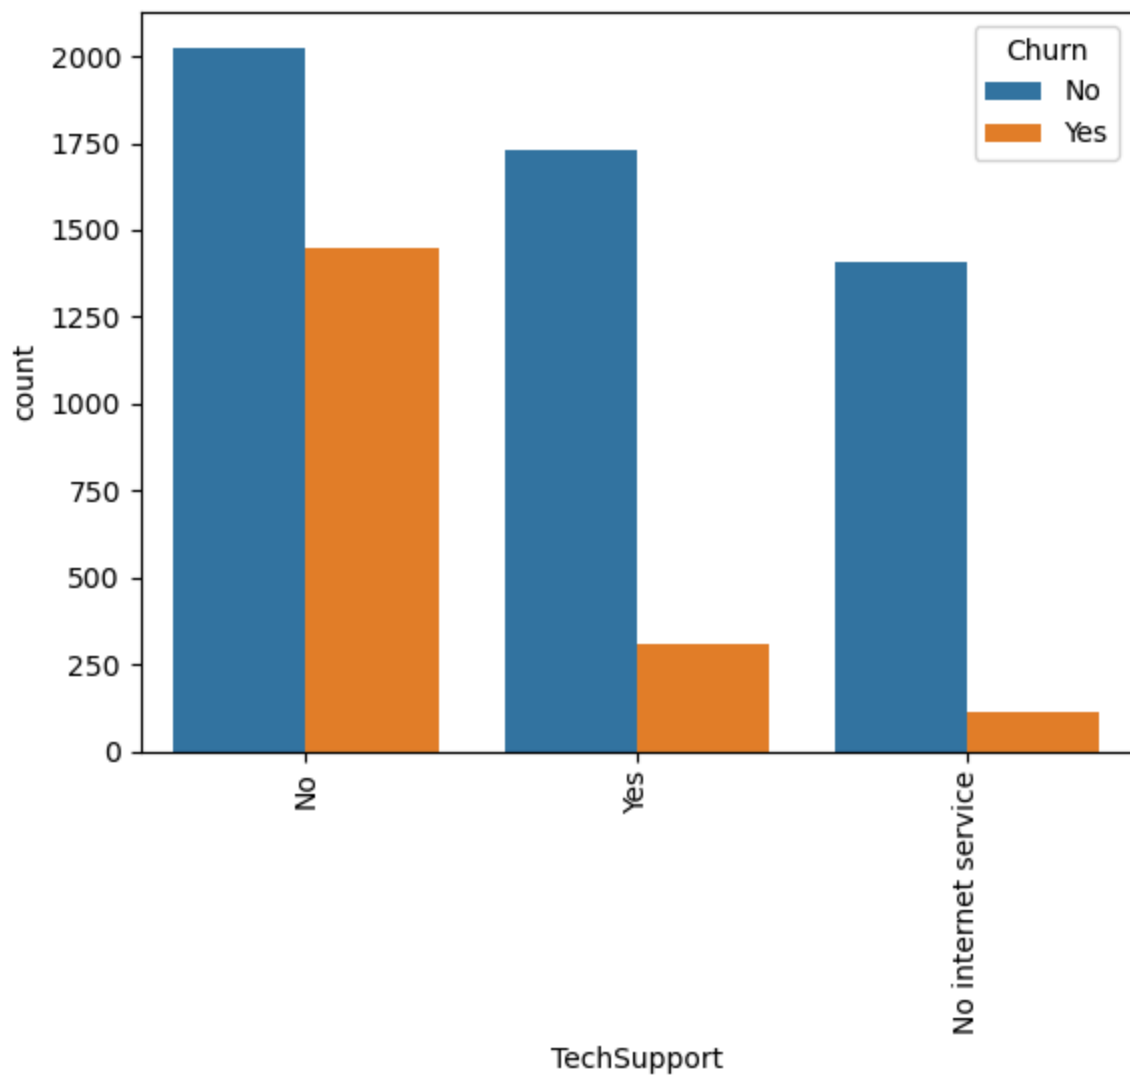


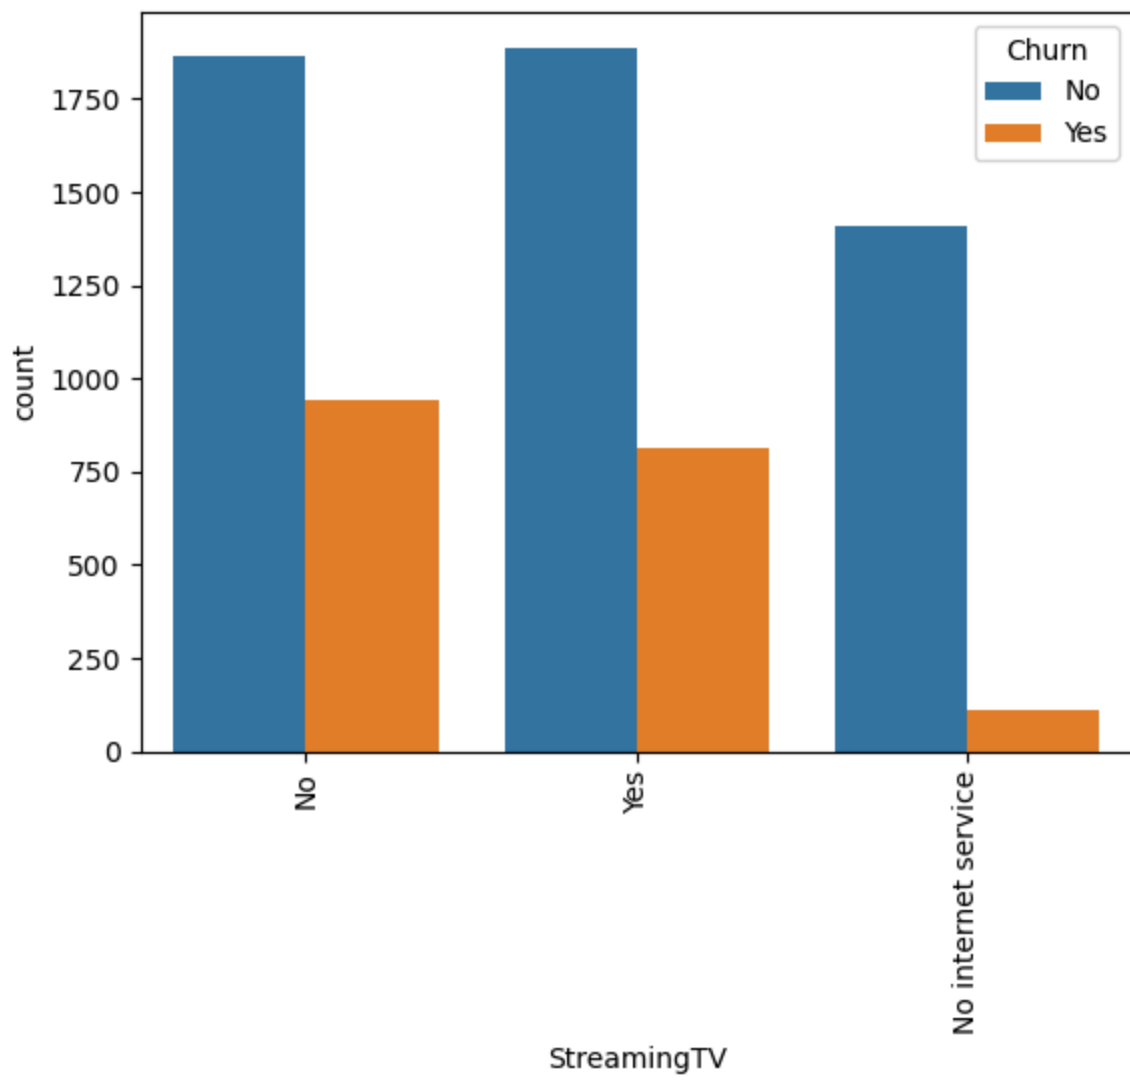





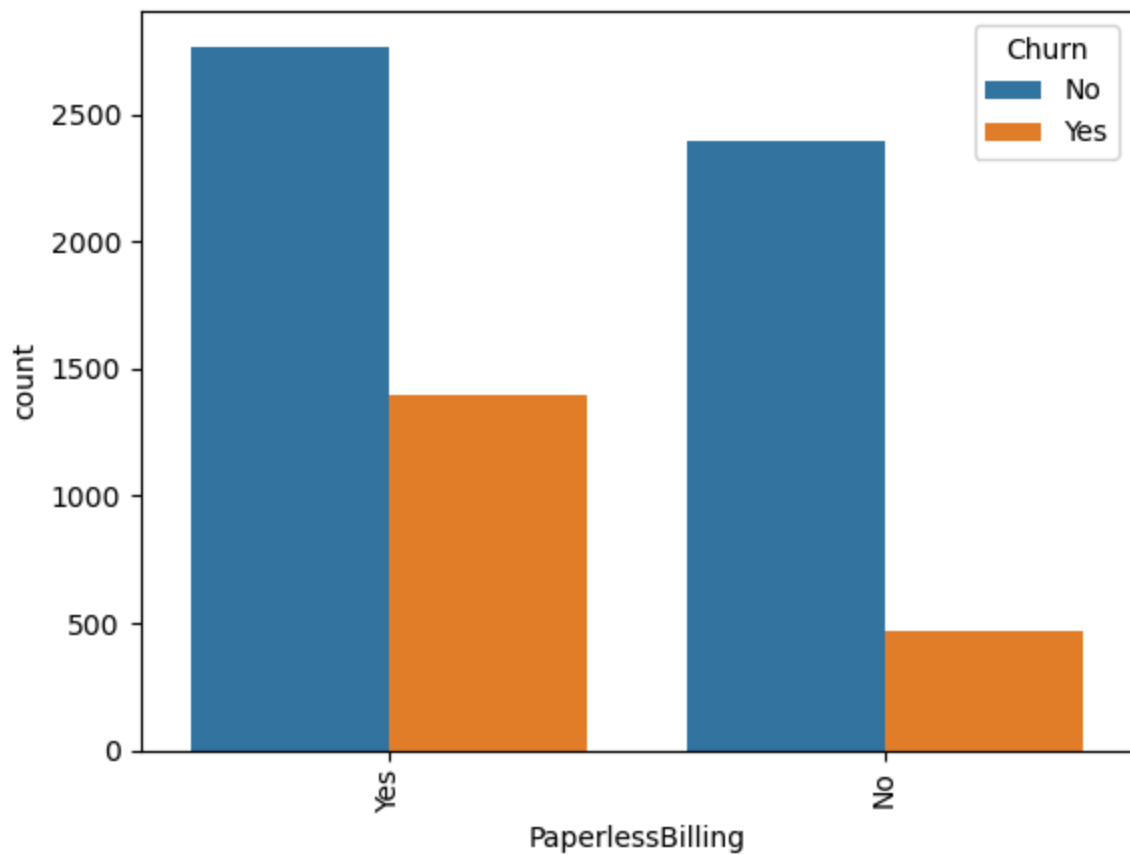



















*convert the target variable into No=0 and Yes=1.

```
In [23]: data["Churn"]=np.where(data.Churn=="Yes",1,0)
```

```
In [24]: data.head()
```

```
Out[24]:
```

	gender	SeniorCitizen	Partner	Dependents	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineSecurity
0	Female	0	Yes	No	No	No phone service	DSL	No	
1	Male	0	No	No	Yes	No	DSL	Yes	
2	Male	0	No	No	Yes	No	DSL	Yes	
3	Male	0	No	No	No	No phone service	DSL	Yes	
4	Female	0	No	No	Yes	No	Fiber optic	No	

```
In [25]: # converting all categorical variabls into dummy variabls.
data_dummies=pd.get_dummies(data)
data_dummies.head()
```

Out[25]:


	SeniorCitizen	MonthlyCharges	TotalCharges	Churn	gender_Female	gender_Male	Partner_No	Partner_Yes
0	0	29.85	29.85	0	1	0	0	1
1	0	56.95	1889.50	0	0	1	1	0
2	0	53.85	108.15	1	0	1	1	0
3	0	42.30	1840.75	0	0	1	1	0
4	0	70.70	151.65	1	1	0	1	0

5 rows × 9 columns

In [26]: `# Relationship between monthly charges and total charges.`

In [27]: `sns.lmplot(data=data_dummies,x="MonthlyCharges",y="TotalCharges",fit_reg=False)`


Out[27]: `<seaborn.axisgrid.FacetGrid at 0x2abd5e35f10>`



- Total charges increase as Monthly charges increases.
- churn by monthly charges and total charges.

In [28]: `Mth=sns.kdeplot(data_dummies.MonthlyCharges[(data_dummies["Churn"]==0)],color="blue",sha
Mth=sns.kdeplot(data_dummies.MonthlyCharges[(data_dummies["Churn"]==1)],ax=Mth,color="re
Mth.legend(["No churn","churn"],loc="upper right")
Mth.set_ylabel("Density")
Mth.set_xlabel("Monthly charges")
Mth.set_title("Monthly charges by churn")`


Out[28]: Text(0.5, 1.0, 'Monthly charges by churn')



- churn is high when monthly charges are high.

```
In [29]: Tot=sns.kdeplot(data_dummies.TotalCharges[(data_dummies["Churn"]==0)],color="blue",shade
Tot=sns.kdeplot(data_dummies.TotalCharges[(data_dummies["Churn"]==1)],ax=Tot,color="red"
Tot.legend(["No churn","churn"],loc="upper right")
Tot.set_ylabel("Density")
Tot.set_xlabel("Total charges")
Tot.set_title("Total charges by churn")
```


Out[29]: Text(0.5, 1.0, 'Total charges by churn')



- High churn at lower total charges.

```
In [30]: # Build a correlation of all predictors with churn.  
plt.figure(figsize=(14,6))  
data_dummies.corr()["Churn"].sort_values(ascending=False).plot(kind="bar")
```

```
Out[30]: <Axes: >
```




Insights.

- Higher rate of churn can be seen in month_to_month contract, No_online security, No technical support, Fibre_optic internet service.
- Lower rate of churn can be seen in Two year contract, subscription without internet service and the customers engaged for 5+ years.
- Gender and phone service has no impact on churn.

```
In [31]: # This is also evident from Heatmap.
plt.figure(figsize=(14,12),dpi=130)
sns.heatmap(data_dummies.corr(),cmap="Paired")
```

```
Out[31]: <Axes: >
```

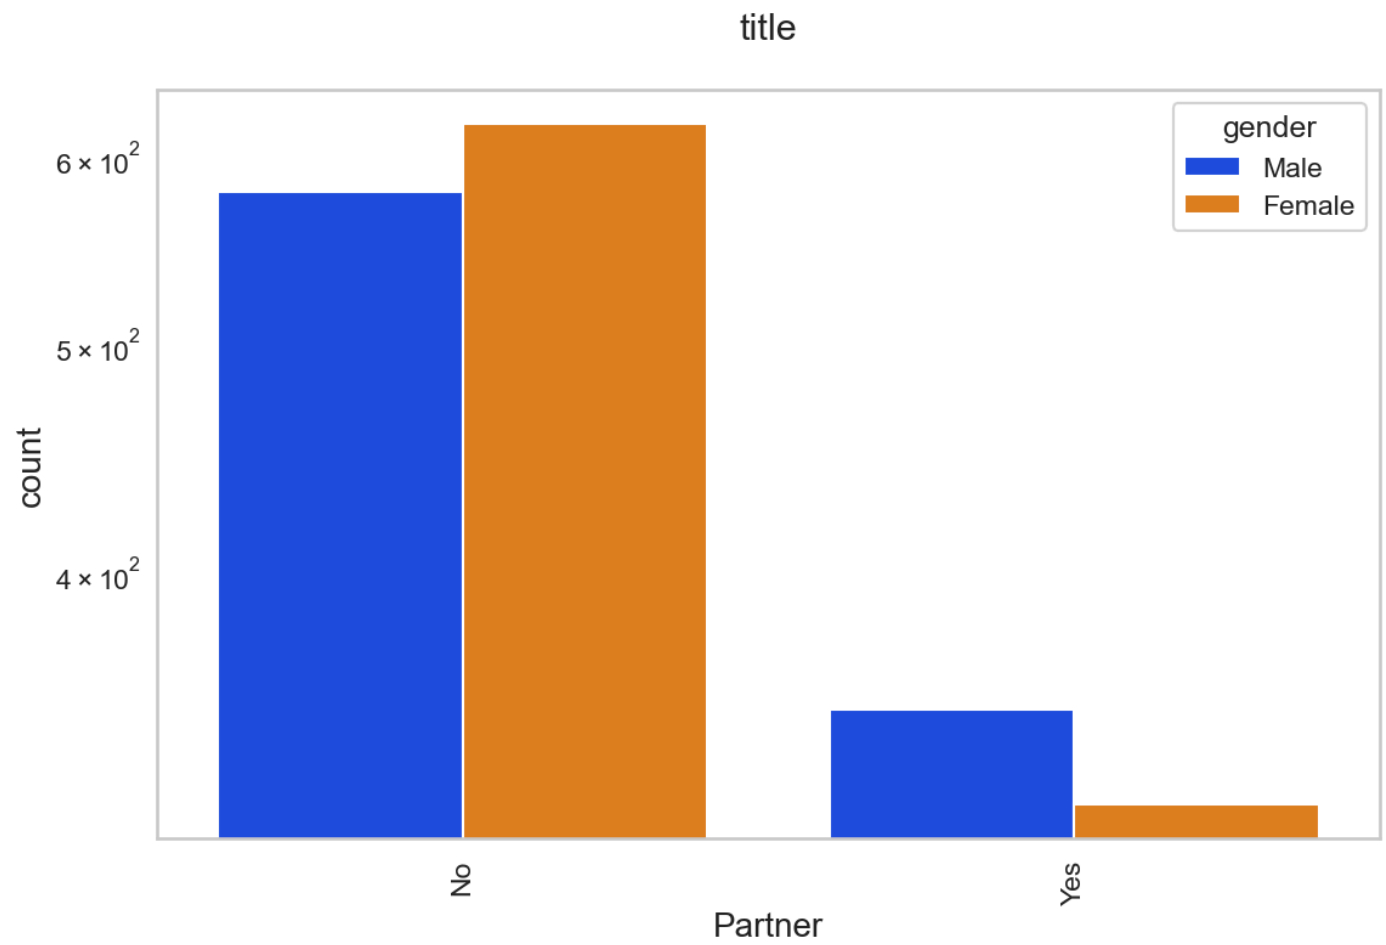
Bivariate Analysis.

```
In [32]: new_data_target0=data.loc[data["Churn"]==0]
new_data_target1=data.loc[data["Churn"]==1]
```


```
In [33]: def uniplot(df,col,title,hue=None):
    sns.set_style("whitegrid")
    sns.set_context("talk")
    plt.rcParams["axes.labelsize"]=20
    plt.rcParams["axes.titlesize"]=22
    plt.rcParams["axes.titlepad"]=30

    temp=pd.Series(data=hue)
    fig,ax=plt.subplots()
    width=len(df[col].unique()+7+4*len(temp.unique()))
    fig.set_size_inches(width,8)
    plt.xticks(rotation=90)
    plt.yscale("log")
    plt.title("title")
    ax=sns.countplot(data=df,x=col,order=df[col].value_counts().index,hue=hue,palette="b
    plt.show
```


```
In [34]: uniplot(new_data_target1,col="Partner",title="distribution of gender for churned custome
```




```
In [35]: uniplot(new_data_target1,col="PaymentMethod",title="distribution of gender for churned c
```



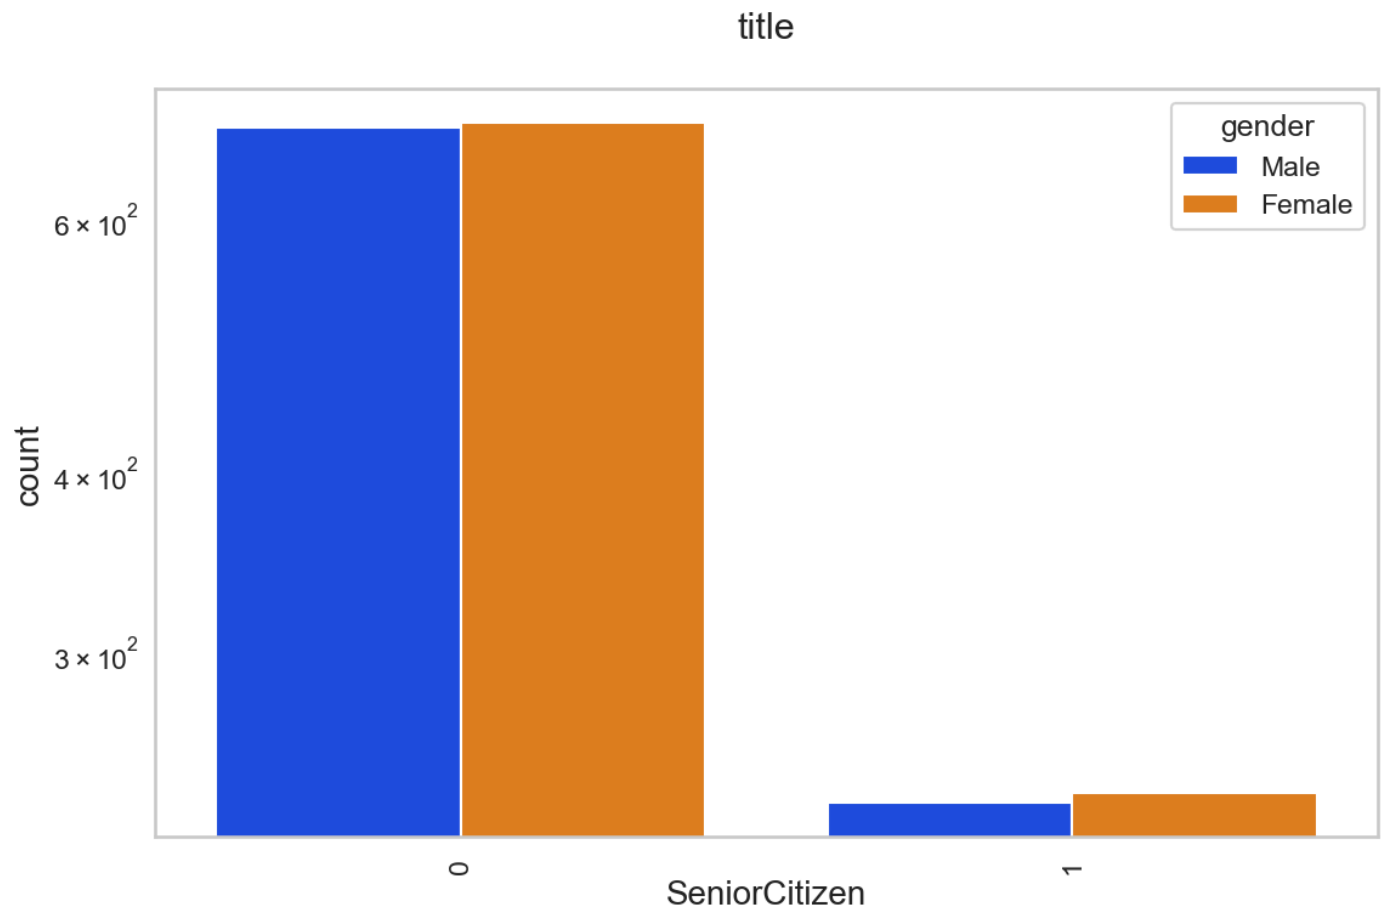
```
In [36]: uniplot(new_data_target1,col="Contract",title="distribution of gender for churned custom
```



```
In [37]: unipLOT(new_data_target1,col="TechSupport",title="distribution of gender for churned cus
```



```
In [38]: uniplot(new_data_target1,col="SeniorCitizen",title="distribution of gender for churned c
```



Conclusions

- Electronic check medium has the highest number of churn customers.
- Monthly customers are more likely to churn as there is no contract.
- Non senior citizens are high in number in terms of churn.
- No online security, No tech support category are high churners.

```
In [39]: data_dummies.to_csv("data.csv")
```

```
In [40]: import pandas as pd
from sklearn import metrics
from sklearn.model_selection import train_test_split
from sklearn.metrics import recall_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.tree import DecisionTreeClassifier
#from imblearn.combine import SMOTEENN
```

```
In [41]: df=pd.read_csv("data.csv")
```

```
In [42]: df.head()
```

Out[42]:

Unnamed: 0

	SeniorCitizen	MonthlyCharges	TotalCharges	Churn	gender_Female	gender_Male	Partner_No	P
0	0	0	29.85	29.85	0	1	0	0
1	1	0	56.95	1889.50	0	0	1	1
2	2	0	53.85	108.15	1	0	1	1
3	3	0	42.30	1840.75	0	0	1	1
4	4	0	70.70	151.65	1	1	0	1

5 rows × 52 columns

In [43]:

df=df.drop("Unnamed: 0",axis=1)

In [44]:

df.head()

Out[44]:

	SeniorCitizen	MonthlyCharges	TotalCharges	Churn	gender_Female	gender_Male	Partner_No	Partner_Yes
0	0	29.85	29.85	0	1	0	0	1
1	0	56.95	1889.50	0	0	1	1	0
2	0	53.85	108.15	1	0	1	1	0
3	0	42.30	1840.75	0	0	1	1	0
4	0	70.70	151.65	1	1	0	1	0

5 rows × 51 columns

In [45]:

creating x and y variables.

x=df.drop("Churn",axis=1)

y=df["Churn"]

In [46]:

x

Out[46]:

	SeniorCitizen	MonthlyCharges	TotalCharges	gender_Female	gender_Male	Partner_No	Partner_Yes	Dep
0	0	29.85	29.85	1	0	0	1	
1	0	56.95	1889.50	0	1	1	0	
2	0	53.85	108.15	0	1	1	0	
3	0	42.30	1840.75	0	1	1	0	
4	0	70.70	151.65	1	0	1	0	
...
7027	0	84.80	1990.50	0	1	0	1	
7028	0	103.20	7362.90	1	0	0	1	
7029	0	29.60	346.45	1	0	0	1	
7030	1	74.40	306.60	0	1	0	1	
7031	0	105.65	6844.50	0	1	1	0	

7032 rows × 50 columns

```
Out[47]: 0 0
          1 0
          2 1
          3 0
          4 1
          ..
          7027 0
          7028 0
          7029 0
          7030 1
          7031 0
          Name: Churn, Length: 7032, dtype: int64
```

```
In [48]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Decision Tree Classifier.

```
In [49]: model_dt=DecisionTreeClassifier(criterion="gini",random_state=42,max_depth=6,min_samples
```

```
In [50]: model_dt.fit(x_train,y_train)
```

```
Out[50]: ▼ DecisionTreeClassifier
          DecisionTreeClassifier(max_depth=6, min_samples_leaf=8, random_state=42)
```

```
In [51]: y_pred=model_dt.predict(x_test)
```

```
In [52]: y_pred
```

```
Out[52]: array([0, 0, 0, ..., 0, 1, 0], dtype=int64)
```

```
In [53]: print(classification_report(y_test,y_pred,labels=[0,1]))
```

	precision	recall	f1-score	support
0	0.84	0.89	0.87	1528
1	0.66	0.56	0.61	582
accuracy			0.80	2110
macro avg	0.75	0.72	0.74	2110
weighted avg	0.79	0.80	0.79	2110

```
In [59]: print(confusion_matrix(y_test,y_pred))
```

```
[[1190  359]
 [ 428  133]]
```

```
In [60]: from sklearn.model_selection import train_test_split
          from sklearn.linear_model import LogisticRegression
          from sklearn.metrics import classification_report

          # Assuming you have your data stored in X (features) and y (target)
          # Split the data into training and testing sets
          x_train, y_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=42

          # Initialize the logistic regression model
          model = LogisticRegression()

          # Train the model on the training data
          model.fit(x_train, y_train)
```



```
# Make predictions on the test data
```

```
predictions = model.predict(x_test)
```

```
# Evaluate the model
```

```
print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.73	0.78	0.76	1549
1	0.26	0.21	0.23	561
accuracy			0.63	2110
macro avg	0.50	0.50	0.49	2110
weighted avg	0.61	0.63	0.62	2110

In []: